

tornikeo_e3

January 31, 2024

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import matplotlib as mpl
import pandas as pd
from scipy import integrate
```

1 Problem 1

Consider the SIR model, consider modelling of synthetic data, and then make parameter estimation for β and γ .

- Get either an analytical solver, or a numerical solver, which outputs $S(t)$, $I(t)$, $R(t)$.

```
[23]: # Adapded from https://scientific-python.readthedocs.io/en/latest/notebooks_rst/
      ↪ 3_Ordinary_Differential_Equations/02_Examples/Epidemic_model_SIR.html

N = 100.
IO, RO = 1.,0
SO = N - IO - RO

beta, gamma = .2, .05

tmax = 100
Nt = tmax

t = np.linspace(0, tmax, Nt + 1)

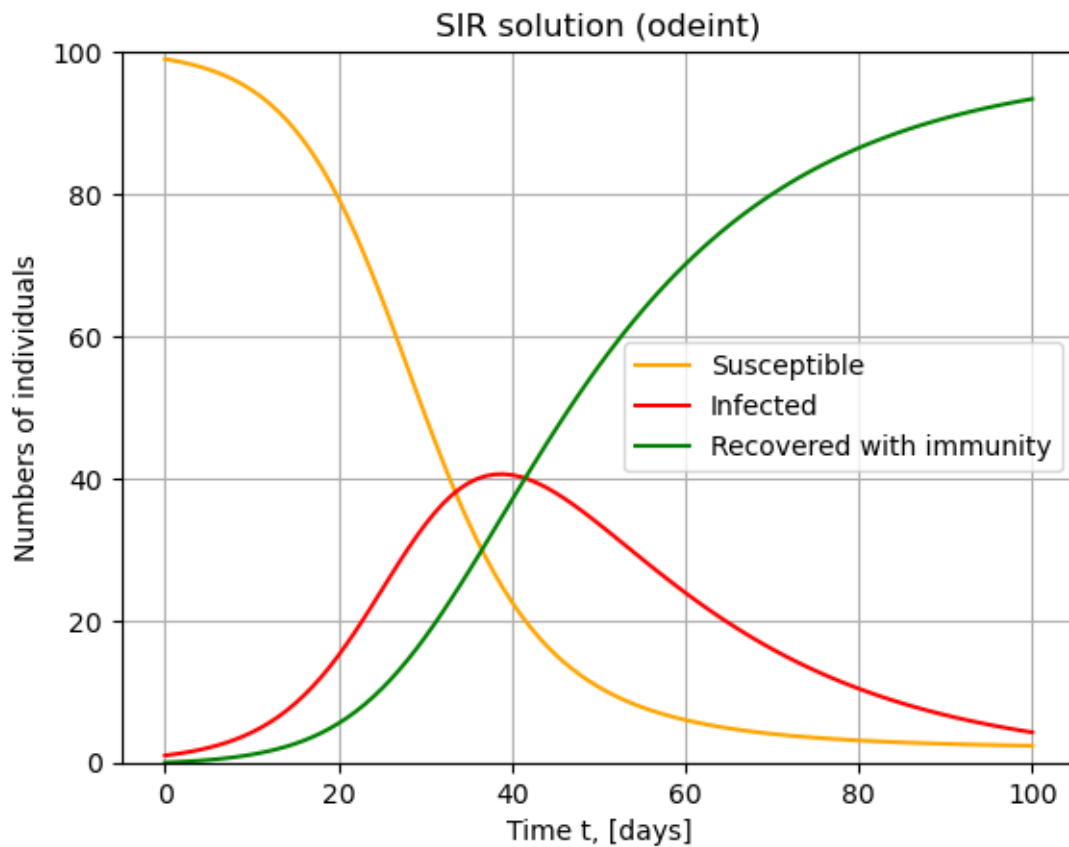
def sir_model(y, t, beta, gamma):
    S, I, R = y
    dSdt = -beta * I * S / N
    dIdt = beta * I * S / N - gamma * I
    dRdt = gamma * I
    return [dSdt, dIdt, dRdt]

res = integrate.odeint(sir_model, (SO, IO, RO), t, args=(beta, gamma))
S,I,R = res.T
Seuil = 1 - 1/(beta/gamma)
print(Seuil)
```

```
plt.figure()
plt.grid()
plt.title("SIR solution (odeint)")
plt.plot(t, S, 'orange', label='Susceptible')
plt.plot(t, I, 'r', label='Infected')
plt.plot(t, R, 'g', label='Recovered with immunity')
plt.xlabel('Time t, [days]')
plt.ylabel('Numbers of individuals')
plt.ylim([0,N])
plt.legend()

plt.show();
```

0.75



- Make synthetic test case with noise-perturbed observations, e.g. $S_{observed}(t) = S_{truth}(t) + e(1)$ where $e \sim N(0, 2I)$

```
[24]: import pandas as pd
from scipy import stats
```

```

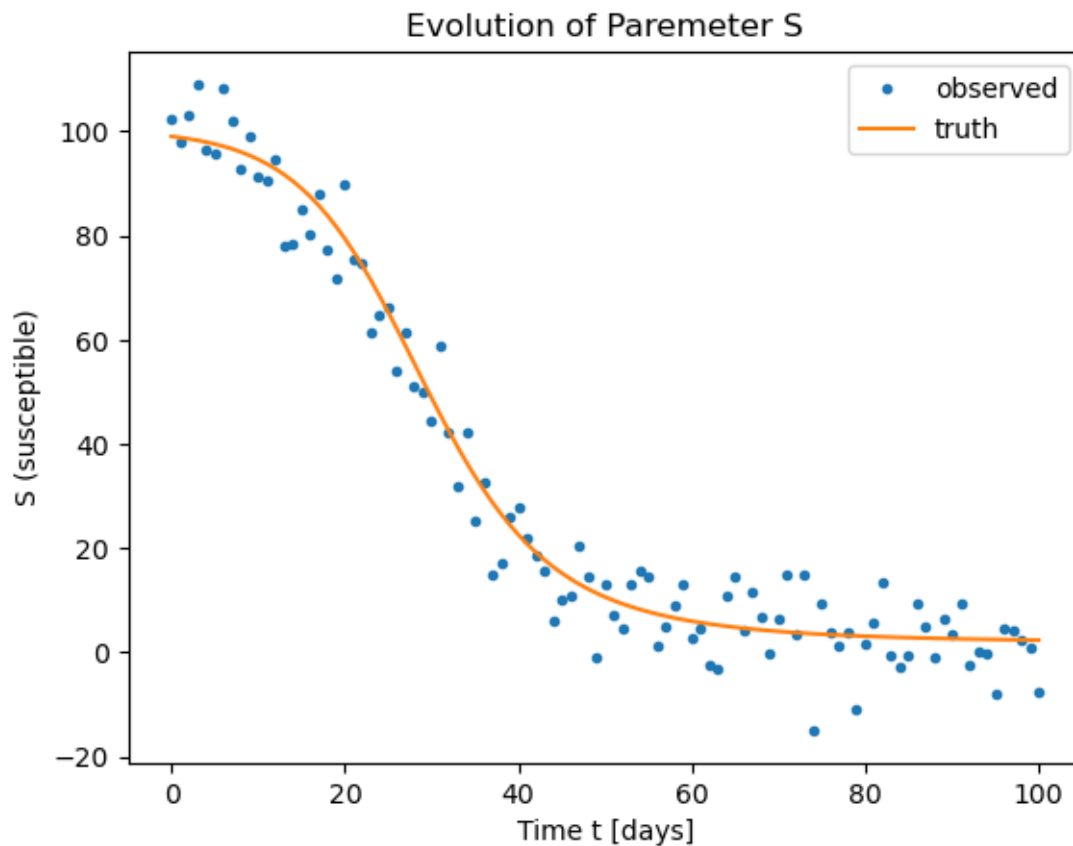
np.random.seed(42)

err_sigma = (N * .5) ** .5
error_fn = stats.norm(0, err_sigma)
noise = pd.Series(error_fn.rvs(size=S.shape))
Sobs = pd.Series(S + noise)

plt.figure()
Sobs.plot(style='.',label='observed')
pd.Series(S).plot(label='truth')
plt.legend()
plt.title("Evolution of Paremeter S")
plt.ylabel('S (susceptible)')
plt.xlabel("Time t [days]")

```

[24]: Text(0.5, 0, 'Time t [days]')



- Write posterior distribution for (β, γ) .

Define SIR model:

$$\dot{S} = -\beta/NI(t)S(t)\dot{I} = \beta/NI(t)S(t) - \gamma I(t)\dot{R} = \gamma I(t)$$

Annotate

$$\theta = [\beta, \gamma]^T$$

Start from bayes theorem:

$$P(\theta|S(t)) \propto P(S(t)|\theta)P(\theta)$$

Which expands to

$$= P(\theta) \frac{1}{\sqrt{((2\pi)^n |\Sigma|)}} \exp(-1/2 S(t)^T \Sigma^{-1} S(t))$$

- Use MCMC to get posterior estimates and uncertainty quantification, and evaluate MCMC chains.

```
[41]: import numpy as np
import pandas as pd
from scipy.integrate import odeint
import matplotlib.pyplot as plt

# Metropolis-Hastings MCMC algorithm
def metropolis_hastings_mcmc(Sobs, iterations, proposal_std):
    # Initial parameter values
    beta_current, gamma_current = 0.1, 0.1

    # Number of data points
    N = len(Sobs)

    # Proposal function (normal distribution)
    def propose_parameters(beta, gamma):
        beta_proposed = stats.norm(beta, proposal_std).rvs()
        gamma_proposed = stats.norm(gamma, proposal_std).rvs()
        return beta_proposed, gamma_proposed

    # Prior function (assuming flat priors)
    def prior(beta, gamma):
        return 1.0

    # Log likelihood function
    def log_likelihood(Y, X, sigma):
        return -0.5 / sigma**2 * np.linalg.norm(Y - X)**2

    # Run MCMC
    beta_chain, gamma_chain = [], []
    t = np.linspace(0, tmax, Nt + 1)
```

```

sigma = err_sigma

for _ in range(iterations):
    # Simulate SIR model for current parameters
    sir_solution_current = odeint(sir_model, [S0, I0, R0], t,
    ↪args=(beta_current, gamma_current))
    S_current = sir_solution_current[:, 0]

    # Propose new parameters
    beta_proposed, gamma_proposed = propose_parameters(beta_current,
    ↪gamma_current)

    # Simulate SIR model for proposed parameters
    sir_solution_proposed = odeint(sir_model, [S0, I0, R0], t,
    ↪args=(beta_proposed, gamma_proposed))
    S_proposed = sir_solution_proposed[:, 0]

    # Compute prior, log likelihood, and posterior
    prior_current = prior(beta_current, gamma_current)
    prior_proposed = prior(beta_proposed, gamma_proposed)
    log_likelihood_current = log_likelihood(Sobs, S_current, sigma)
    log_likelihood_proposed = log_likelihood(Sobs, S_proposed, sigma)

    log_posterior_current = np.log(prior_current) + log_likelihood_current
    log_posterior_proposed = np.log(prior_proposed) +
    ↪log_likelihood_proposed

    # Accept or reject the proposal
    acceptance_ratio = np.exp(log_posterior_proposed -
    ↪log_posterior_current)
    if np.random.uniform(0, .3) < acceptance_ratio:
        beta_current, gamma_current = beta_proposed, gamma_proposed

    # Save current parameters to chains
    beta_chain.append(beta_current)
    gamma_chain.append(gamma_current)

return np.array(beta_chain), np.array(gamma_chain)

# Example usage
# Assuming Sobs is the measured data (pandas series)
N = len(Sobs)
S0, I0, R0 = int(Sobs.iloc[0]), 1, 0 # Initial conditions

# Set the standard deviation for the proposal distribution
proposal_std = .05

```

```

# Number of MCMC iterations
iterations = 8000

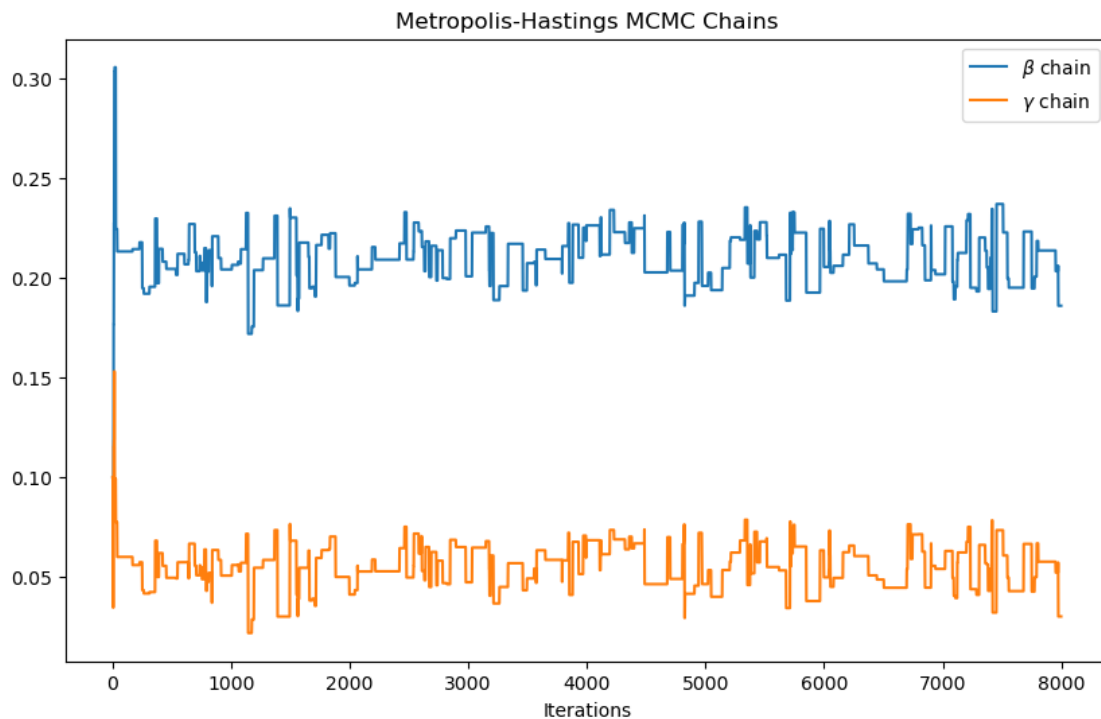
# Run MCMC
beta_chain, gamma_chain = metropolis_hastings_mcmc(Sobs.values, iterations,
↪proposal_std)

# Plot MCMC chains
plt.figure(figsize=(10, 6))
plt.plot(beta_chain, label=r'\beta$ chain')
plt.plot(gamma_chain, label=r'\gamma$ chain')
plt.legend()
plt.title('Metropolis-Hastings MCMC Chains')
plt.xlabel('Iterations')
plt.show()

```

/tmp/ipykernel_8600/2182796702.py:56: RuntimeWarning: overflow encountered in exp

```
acceptance_ratio = np.exp(log_posterior_proposed - log_posterior_current)
```



```

[42]: # Discard burn-in period
beta_chain = beta_chain[500:]
gamma_chain = gamma_chain[500:]

```

```

[43]: # Plot trace plots
plt.figure(figsize=(12, 8))

plt.subplot(2, 1, 1)
plt.plot(beta_chain, label=r'$\beta$ chain')
plt.axhline(np.mean(beta_chain), color='red', linestyle='dashed', label='Mean')
plt.title(r'Trace Plot for $\beta$')
plt.xlabel('Iterations')
plt.legend()

plt.subplot(2, 1, 2)
plt.plot(gamma_chain, label=r'$\gamma$ chain')
plt.axhline(np.mean(gamma_chain), color='red', linestyle='dashed', label='Mean')
plt.title(r'Trace Plot for $\gamma$')
plt.xlabel('Iterations')
plt.legend()

plt.tight_layout()
plt.show()

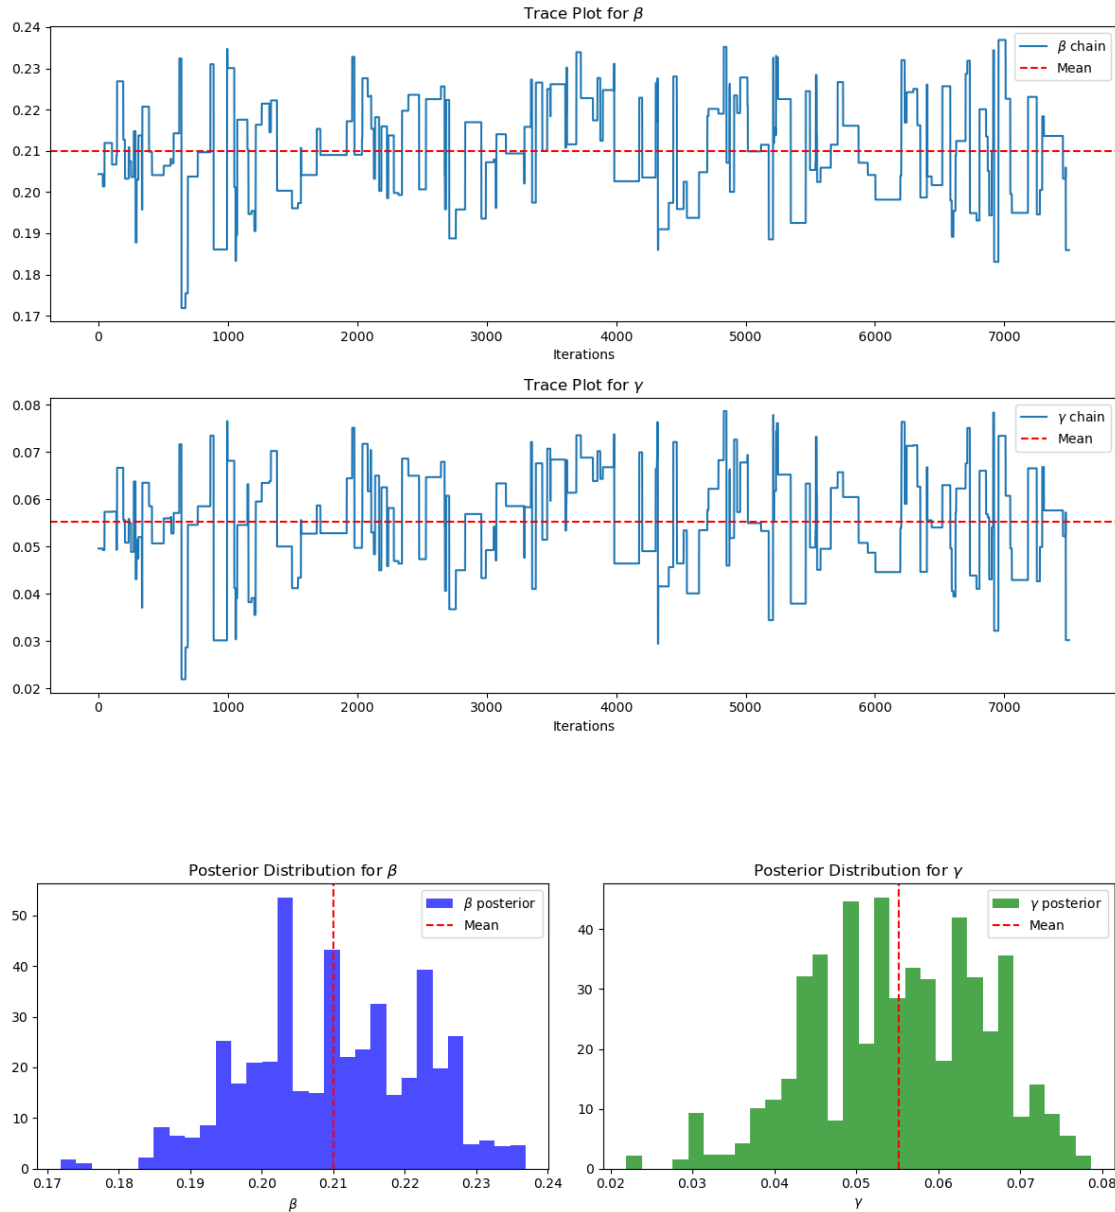
# Plot posterior distributions
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.hist(beta_chain, bins=30, density=True, alpha=0.7, color='blue',
         label=r'$\beta$ posterior')
plt.axvline(np.mean(beta_chain), color='red', linestyle='dashed', label='Mean')
plt.title(r'Posterior Distribution for $\beta$')
plt.xlabel(r'$\beta$')
plt.legend()

plt.subplot(1, 2, 2)
plt.hist(gamma_chain, bins=30, density=True, alpha=0.7, color='green',
         label=r'$\gamma$ posterior')
plt.axvline(np.mean(gamma_chain), color='red', linestyle='dashed', label='Mean')
plt.title(r'Posterior Distribution for $\gamma$')
plt.xlabel(r'$\gamma$')
plt.legend()

plt.tight_layout()
plt.show()

```



- Plot predictive intervals.

```
[45]: # Function to simulate multiple trajectories
def simulate_trajectories(parameters, num_trajectories, t):
    trajectories = []
    for i in range(num_trajectories):
        beta, gamma = parameters[i]
        sir_solution = odeint(sir_model, [S0, I0, R0], t, args=(beta, gamma))
        S_trajectory = sir_solution[:, 0]
        trajectories.append(S_trajectory)
```



```

    return np.array(trajectories)

# Number of trajectories to simulate
num_trajectories = 1000

# Simulate trajectories using MCMC samples
parameters_samples = np.column_stack((beta_chain[100:], gamma_chain[100:]))
trajectories = simulate_trajectories(parameters_samples, num_trajectories, t)

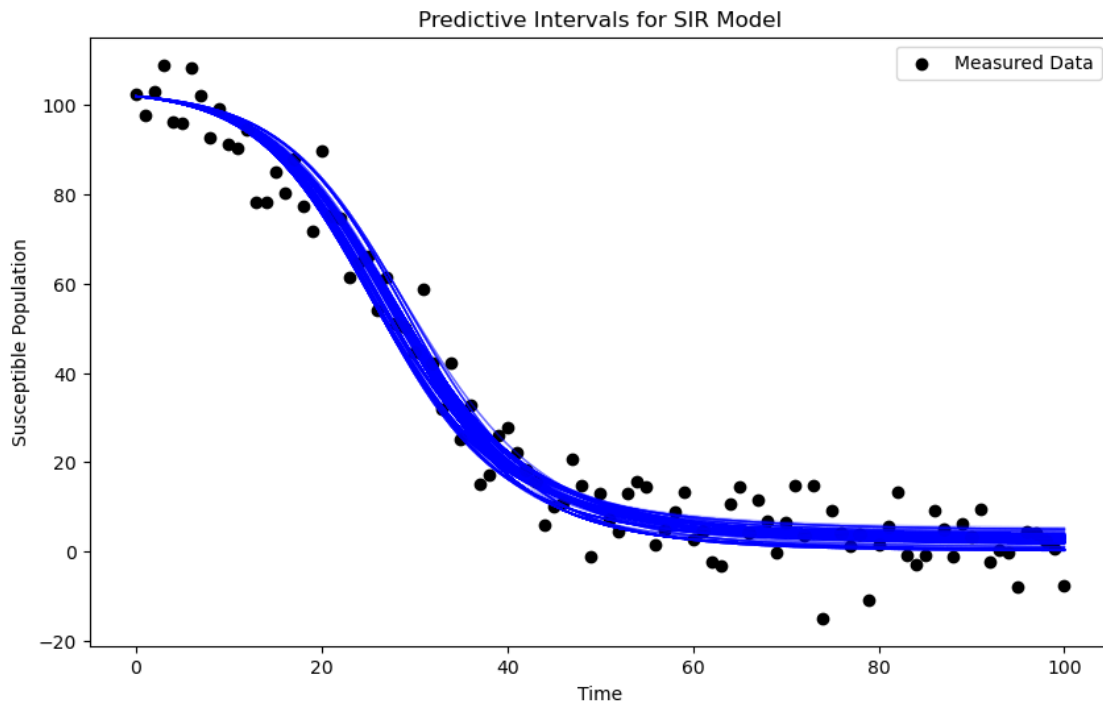
# Plot predictive intervals
plt.figure(figsize=(10, 6))

# Plot measured data
plt.scatter(t, Sobs, color='black', marker='o', label='Measured Data')

# Plot simulated trajectories
for i in range(num_trajectories):
    plt.plot(t, trajectories[i], color='blue', alpha=0.1, linewidth=1)

plt.title('Predictive Intervals for SIR Model')
plt.xlabel('Time')
plt.ylabel('Susceptible Population')
plt.legend()
plt.show()

```



1.1 Problem 2

Take one realisation of white noise, form posterior distribution, and obtain variance estimate with MCMC. Make the standard plots.

Generate white noise.

```
[46]: import numpy as np
```

```
np.random.seed(42) # Set seed for reproducibility
white_noise = np.random.normal(0, 1, size=100) # Adjust size as needed
```

```
[47]: import pymc as pm
```

```
with pm.Model() as model:
    sigma = pm.HalfNormal('sigma', sigma=1) # Prior for the standard deviation,
    ↪(positive values only)
    obs = pm.Normal('obs', mu=0, sigma=sigma, observed=white_noise)
```

```
[48]: with model:
```

```
    trace = pm.sample(1000, tune=1000, cores=4) # Adjust the number of samples,
    ↪and tuning steps
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [sigma]


<IPython.core.display.HTML object>

<IPython.core.display.HTML object>


Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 1 seconds.

```
[49]: pm.plot_trace(trace)
```

```
[49]: array([[<Axes: title={'center': 'sigma'}>,
    <Axes: title={'center': 'sigma'}>]], dtype=object)
```



```
[50]: pm.plot_autocorr(trace);
```



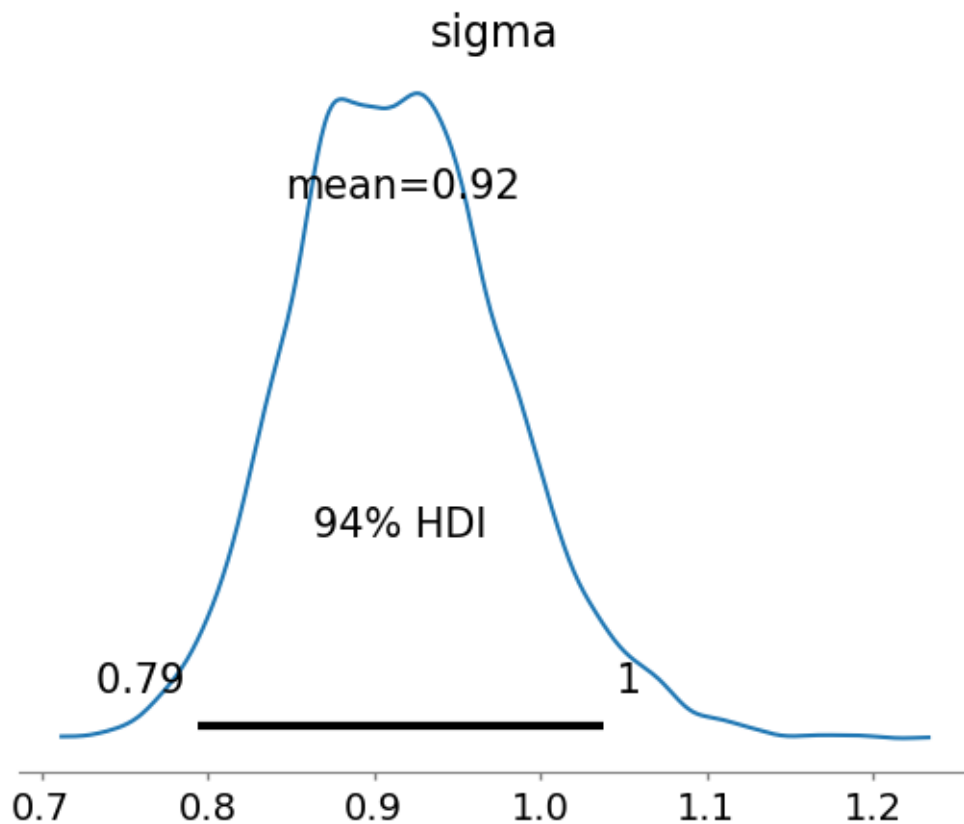
Posterior analysis:

```
[51]: posterior_samples = trace['posterior']['sigma']
estimated_variance = np.mean(posterior_samples)
print("Estimated Variance:", estimated_variance)
```

```
Estimated Variance: <xarray.DataArray 'sigma' ()>
array(0.9168637)
```

```
[52]: pm.plot_posterior(trace)
```

```
[52]: <Axes: title={'center': 'sigma'}>
```



1.2 Problem 3

Given one realisation of the Ornstein-Uhlenbeck process with fixed δ , use MCMC to obtain CM-estimate of θ . That is,

- Draw one realisation of the OU process.
- Formulate the posterior of θ .
- Use MCMC to estimate θ .
- Plot chains, ACFs, densities and compute ESS and OES.
- Visualise posterior parameter estimates and Monte Carlo errors

```
[53]: import numpy as np

np.random.seed(42)

def ornstein_uhlenbeck(delta, theta, sigma, x0, n, dt=0.1):
    ou_process = np.zeros(n)
    ou_process[0] = x0
    for t in range(1, n):
        drift = theta * (delta - ou_process[t - 1]) * dt
        diffusion = sigma * np.sqrt(dt) * np.random.normal(0, 1)
        ou_process[t] = ou_process[t - 1] + drift + diffusion
    return ou_process

# Set parameters
delta = 0.0 # Mean-reverting level
theta_true = 1.0 # Mean-reversion parameter (to be estimated)
sigma_true = 0.2 # Volatility parameter
x0 = 0.0 # Initial value
n = 100 # Number of time steps

# Generate OU process realization
ou_data = ornstein_uhlenbeck(delta, theta_true, sigma_true, x0, n)
```

Define model with OU process data.

```
[57]: import pymc as pm

with pm.Model() as model_ou:
    lambda_ = pm.Exponential('lambda_', lam=1) # Prior for the mean-reversion_
    ↪parameter
    obs = pm.Normal('obs', mu=delta + lambda_ * (ou_data[:-1] - delta),
                    sigma=sigma_true, observed=ou_data[1:])
```

```
[58]: # Sample
with model_ou:
```

```
trace_ou = pm.sample(1000, tune=1000, cores=4)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [lambda_]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 1 seconds.

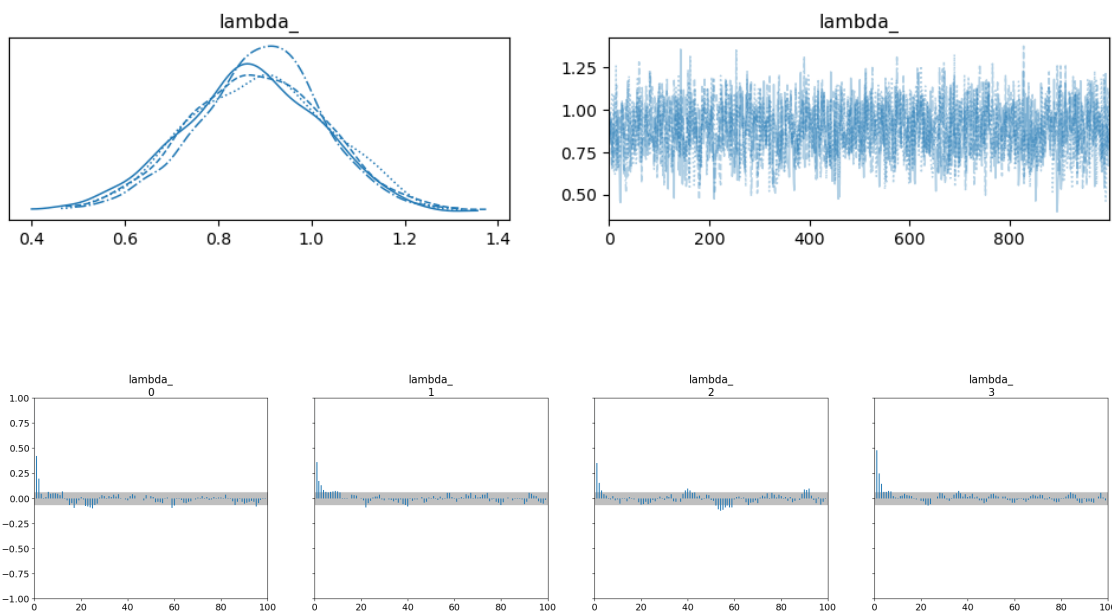
```
[59]: pm.plot_trace(trace_ou)
      pm.plot_autocorr(trace_ou)
      pm.summary(trace_ou).round(2)
```

```
[59]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	\
lambda_	0.88	0.14	0.61	1.16	0.0	0.0	1309.0	1650.0	


```

      r_hat
lambda_    1.0
```

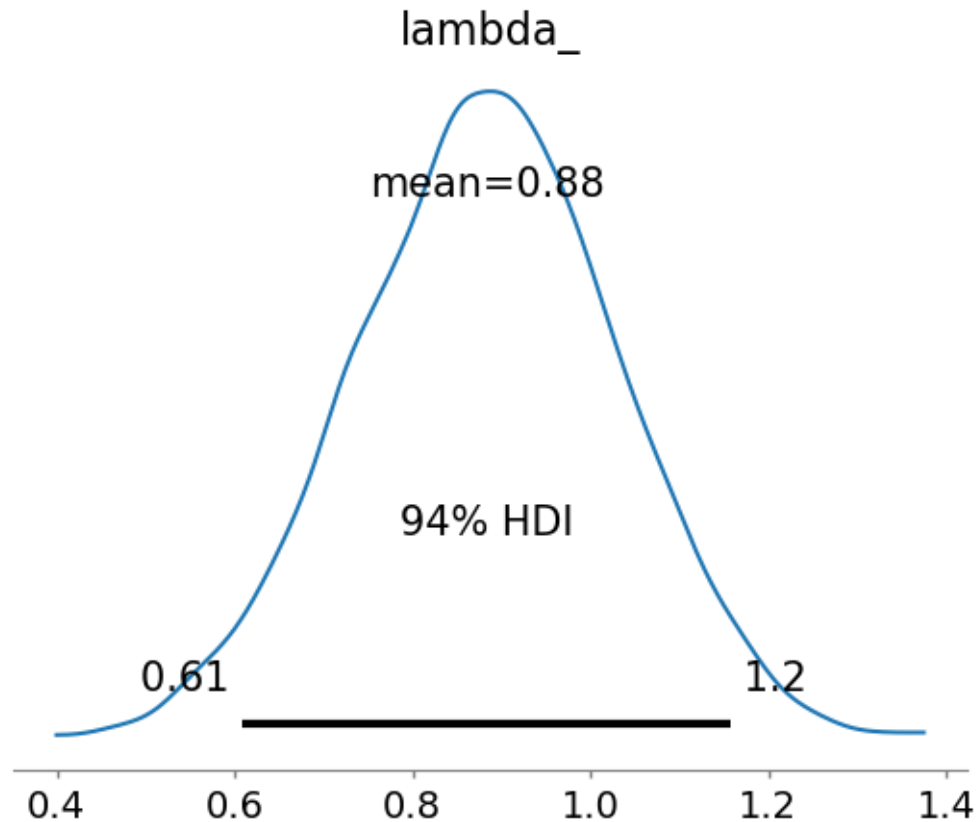


```
[60]: posterior_samples_ou = trace_ou['posterior']['lambda_']
      estimated_lambda = np.mean(posterior_samples_ou)
      print("Estimated Lambda:", estimated_lambda)
```

```
Estimated Lambda: <xarray.DataArray 'lambda_' ()>  
array(0.88323188)
```

```
[66]: pm.plot_posterior(trace_ou);  
  
# Compute ESS and OES  
# ess = pm.ess(trace_ou)  
# oes = pm.oes(trace_ou)  
  
trace_ou.posterior  
# print("Effective Sample Size:", ess)  
# print("OES (Overlapping Estimate of Effective Sample Size):", oes)
```

```
[66]: <xarray.Dataset>  
Dimensions:  (chain: 4, draw: 1000)  
Coordinates:  
    * chain      (chain) int64 0 1 2 3  
    * draw       (draw) int64 0 1 2 3 4 5 6 7 8 ... 992 993 994 995 996 997 998 999  
Data variables:  
    lambda_      (chain, draw) float64 0.8209 0.8209 0.8704 ... 0.671 0.8465 0.8558  
Attributes:  
    created_at:      2024-01-31T18:35:13.088214  
    arviz_version:    0.17.0  
    inference_library: pymc  
    inference_library_version: 5.10.0  
    sampling_time:    1.3804991245269775  
    tuning_steps:     1000
```



2. In order to see different behaviour of priors, do the same as above, but with different priors. Use a Gaussian prior.

```
[72]: import numpy as np
import pymc as pm
import arviz as az

np.random.seed(42)

def ornstein_uhlenbeck(delta, theta, sigma, x0, n, dt=0.1):
    ou_process = np.zeros(n)
    ou_process[0] = x0
    for t in range(1, n):
        drift = theta * (delta - ou_process[t - 1]) * dt
        diffusion = sigma * np.sqrt(dt) * np.random.normal(0, 1)
        ou_process[t] = ou_process[t - 1] + drift + diffusion
    return ou_process

def create_ou_model(prior_lambda):
    with pm.Model() as model_ou:
```

```

        lambda_ = pm.Exponential('lambda_', lam=1) if prior_lambda == '
    ↪ 'exponential' else \
            pm.Normal('lambda_', mu=0, sigma=err_sigma) if prior_lambda
    ↪ == 'gaussian' else \
            pm.Uniform('lambda_') # Flat prior

    obs = pm.Normal('obs', mu=delta + lambda_ * (ou_data[:-1] - delta),
    ↪ sigma=sigma_true, observed=ou_data[1:])

    return model_ou

def report_and_plot(trace, prior_name):
    # Plot chains, ACFs, and densities
    pm.plot_trace(trace)
    pm.plot_autocorr(trace)
    pm.plot_posterior(trace)

    # Compute and print summary statistics
    summary = pm.summary(trace).round(2)
    print("Summary Statistics for", prior_name, "Prior:")
    print(summary)

    # Display the plots
    plt.show()

# Set parameters
delta = 0.0
theta_true = 1.0
sigma_true = 0.2
x0 = 0.0
n = 100

# Generate OU process realization
ou_data = ornstein_uhlenbeck(delta, theta_true, sigma_true, x0, n)

# Priors to compare
prior_names = ['Exponential', 'Gaussian', 'Flat']

for prior_name in prior_names:
    # Create the OU model with the specified prior
    model_ou = create_ou_model(prior_name)

    # MCMC sampling
    with model_ou:
        trace_ou = pm.sample(1000, tune=1000, cores=2)

    # Plotting and reporting

```



```
report_and_plot(trace_ou, prior_name)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

Multiprocess sampling (2 chains in 2 jobs)

NUTS: [lambda_]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>





Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 1 seconds.


We recommend running at least 4 chains for robust computation of convergence diagnostics

Summary Statistics for Exponential Prior:

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	\
lambda_	0.84	0.1	0.66	1.0	0.0	0.0	622.0	611.0	

r_hat
lambda_ 1.0





Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

Multiprocess sampling (2 chains in 2 jobs)

NUTS: [lambda_]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

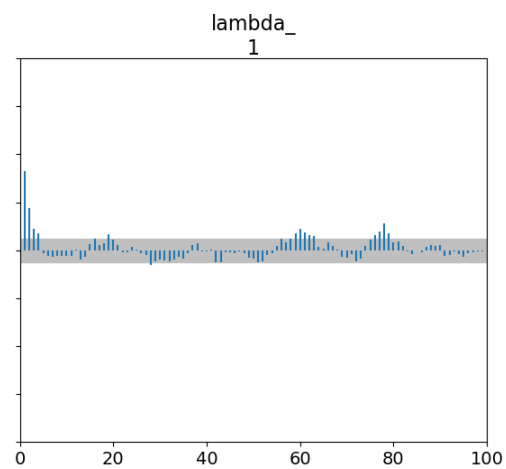



Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 1 seconds.

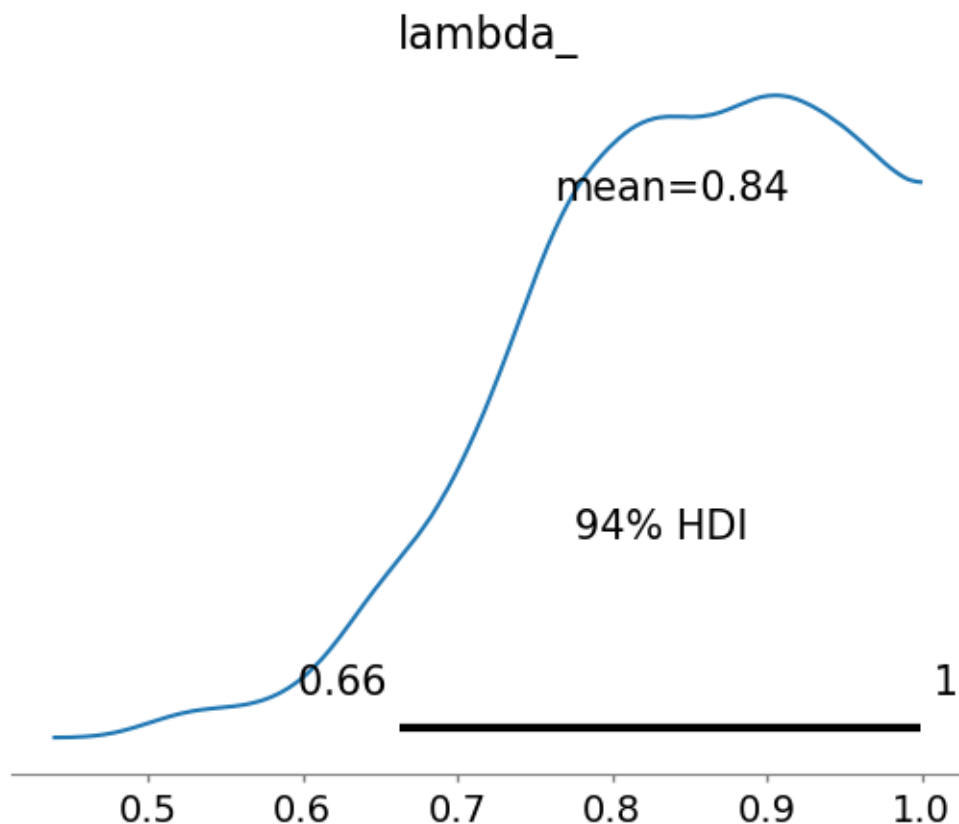
We recommend running at least 4 chains for robust computation of convergence diagnostics

Summary Statistics for Gaussian Prior:

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	\
lambda_	0.84	0.1	0.66	1.0	0.0	0.0	696.0	549.0	

	r_hat
lambda_	1.01





Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

Multiprocess sampling (2 chains in 2 jobs)

NUTS: [lambda_]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>





Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 1 seconds.

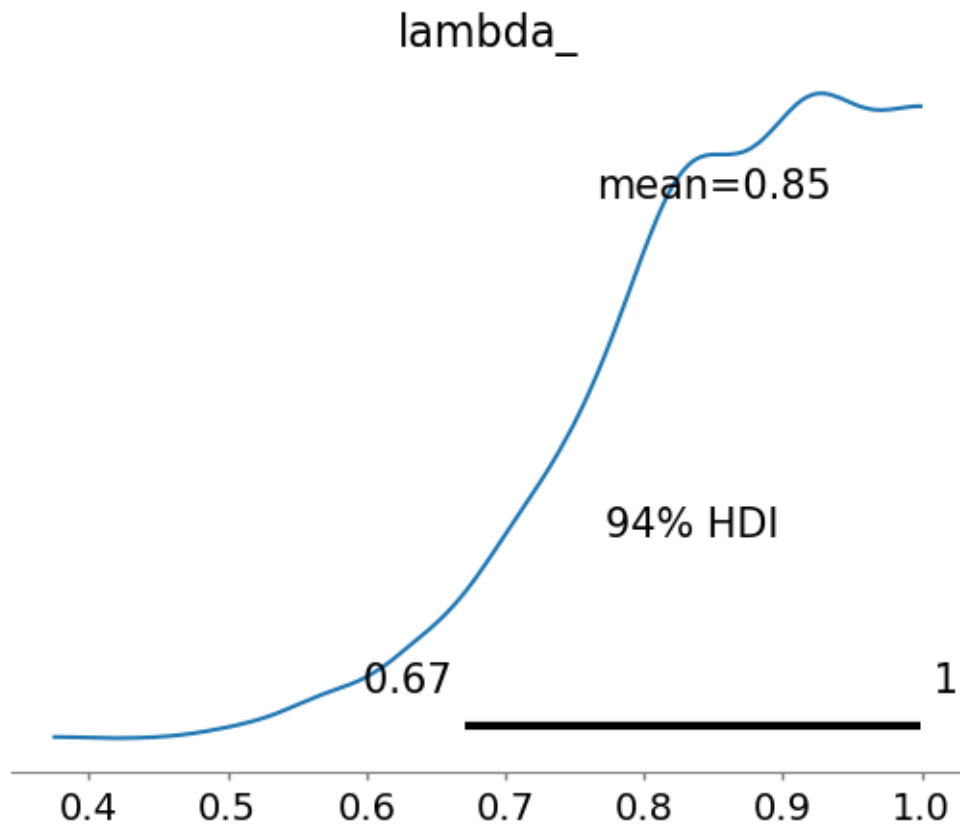
We recommend running at least 4 chains for robust computation of convergence diagnostics

Summary Statistics for Flat Prior:

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	\
lambda_	0.85	0.1	0.67	1.0	0.0	0.0	447.0	535.0	

	r_hat
lambda_	1.0





... and that's it, folks. That's as much as I can do given this timeframe. Till next time.